

PRESENTATION SCRIPT

My First Look at the Data

Worth the Investment? Weekly Stakeholder Update

Raw Analysis: Value-by-Age Curves and Injury

Frequency Patterns

Thursday, September 10, 2026 · Presented by Mahmoud
Script follows the "Data_Analysis_Raw.pptx" deck, slide by
slide

How to use this script

- Each card below matches one slide, in order. The gold number is the slide number, so it is easy to check you are on the right one.
- The main paragraph is what to say out loud. It is written in plain, spoken language, not slide text, so it will not sound like you are reading bullet points.
- The gold "if asked" boxes are optional. Only use them if a stakeholder actually asks a follow up question on that slide.

- You do not need to say every line word for word. Treat this as a safety net, not a script to recite.
- This script uses "I" throughout, matching this deck, since this update is framed as your individual analysis work rather than a team narration.

13

slides

~12

min talk time

9:00

prep, per usual
schedule

10:00

discussion
starts

What this meeting needs to cover

The instructor's own message for this Thursday is specific: no particular tool is required, data cleaning should be done, and the meeting should briefly share first insights about cleaning the data and any assumptions made. Tableau is fine to use but not required.

This deck answers that directly on slides 5 and 6, insights from cleaning, then the assumptions behind them, right after the methodology slide. The value-by-age and injury-frequency content afterward is this project's own roadmap-planned raw analysis, kept in because it is real, finished work and adds substance beyond the minimum ask.

OPENING LINE, BEFORE SLIDE 1

"Good morning, thank you for the time. Today I'm sharing what cleaning the data actually revealed, the assumptions I had to make along the way, and my first real look at the data itself: how player value moves with age, and how often players actually get injured. This will take about twelve minutes, and I'm happy to take questions any time."

1

TITLE SLIDE

My First Look at the Data

SAY

"This is my title slide. I'm still on the Worth the Investment project, player value and injury risk, and this is my weekly stakeholder update for September 10th, my first look at the actual data itself."

SAY

"Before any numbers, here is where this sits on my roadmap. Data retrieval and wrangling is done, that's what I covered last week. This week is raw analysis: value-by-age curves and injury frequency patterns. Fine-tune analysis, the forecast and the risk score, is up next, at a future update."

IF ASKED WHY THIS IS SPLIT INTO TWO SEPARATE MEETINGS

It's how I laid out my own roadmap from the start, in `Project_brainstorm.md`. Week 3 covers raw analysis first, then fine-tune. I kept them as two separate stakeholder updates on purpose, so I can show honest, checkable progress each week instead of holding everything back for one big reveal at the end.

3

SCOPE

What This Update Covers

SAY

"To be precise about scope: today covers value-by-age curves and injury frequency patterns, both grounded in real numbers from my dataset. The Prophet value forecast, the injury-risk score, and the combined recruitment signal are coming at the next update, once I've built on top of what I'm showing today."

4

METHOD

How I Explored the Data

SAY

"Same habit as last week: before drawing any conclusions, I checked the data was solid enough to trust. I connected directly to the two dbt marts, no manual export step. I ran sanity checks first, player counts, average valuation snapshots per player, how many players show zero injuries. And I flagged a known limitation up front, injury counts are capped at 60 for a small number of players, a pagination limit in the source data, not a real injury ceiling."

Three cleaning problems shaped every number that follows

SAY

"Since this is what you specifically asked to hear about, here are the real insights cleaning surfaced. Injury dates needed real parsing, they were stored as plain text with a stray dash for unknowns, and left as text, even sorting by most recent injury was silently wrong. I fixed that, **99.98 percent** parsed cleanly, **107,951 of 107,971**.

Two different join keys were needed. Transfermarkt values and injuries share a real internal player ID, **18,824 of 18,825** overlap, so those join directly. Performance data has no shared ID, so it needed a name-based join instead, a more fragile kind of match.

And matching surfaced a real gap, not a bug. A large share of players had no performance record at all once matched. I didn't assume that was an error, I investigated it directly, more on that a few slides from now."

SAY

"And here are the assumptions behind those fixes, stated plainly. A missing injury date, I treated as unknown, not zero, it becomes a real NULL rather than a guessed date. A player with zero recorded minutes, I assumed reflects real limits in what the performance source covers, not that they never played, so I kept them in the dataset and disclosed the gap instead of quietly dropping them.

The injury count cap at 60, I assumed the true count is unknown above that, not that 60 is a real ceiling, so those players are flagged, not treated as fully counted. And for matching performance data by name, I assumed it was reliable enough after manually reviewing the closest near-matches myself and finding no missed true matches, rather than adding heavier fuzzy-matching logic."

SAY

"This is the same combined dataset from last week: **1,959 players, 32,544** valuation snapshots, spanning more than twenty years. It comes from two dbt marts, one row per player per valuation snapshot for value, and one row per player summarizing injury count, days missed, and playing time."

Market value peaks at 24-26, then falls 93% by 36+

SAY

"The classic scouting assumption is that market value peaks in a player's mid-twenties. I tested that directly. For every valuation snapshot, I worked out the player's age at that point, grouped players into three-year age bands, and took the median market value in each band. That's what this chart shows."

IF ASKED HOW I GOT THESE CHARTS

I pulled the real numbers first, straight from the database, not from a charting tool. I connected to the same DuckDB warehouse dbt builds, queried the value and injury marts with Python and pandas, and grouped the results into bands, by age here, by injury count on the next chart. Only once I had those summary numbers did I build the charts inside the PowerPoint file itself, with a script, so they're real, native PowerPoint charts I can still edit, not pictures pasted in from my notebook.

FINDING 1, INTERPRETATION

9

Elite valuations pull the mean 2.5x above the median

SAY

"Three things stand out. Value rises steadily from the late teens into the early twenties, then climbs sharply. It peaks in the **24 to 26** age bracket, median value **5.0 million euros**, then declines in every bracket that follows. This confirms the standard scouting assumption, and it tells me age belongs as a feature when I build the value forecast next."

IF ASKED WHY THE MEAN, 12.4 MILLION, IS SO MUCH HIGHER THAN THE MEDIAN, 5.0 MILLION

That gap is a real signal, not a mistake. A handful of superstar valuations in that age band, well above 100 million euros, pull the average up a lot, while the median stays anchored to what a typical player in that band is actually worth. I'm treating the median as the more representative number for this chart.

40% of players carry no recorded injury at all

SAY

"The second question: how often do players in this dataset actually get injured? I grouped every player by their recorded injury count, into buckets, from zero up to eleven or more, using the same approach as the value chart. **40 percent** of players, **783 of 1,959**, show zero recorded injuries."

IF ASKED WHY THIS MATTERS FOR RECRUITMENT

It means injury risk is concentrated in a minority of players, not spread evenly. That's exactly the kind of pattern the injury-risk score at the next update needs to be built on, since a single average injury rate across all players would hide this.

SAY

"Two limitations I'm carrying into the fine-tune stage, worth saying out loud now rather than later. The injury count cap: **12 players** sit exactly at 60, so their true injury count beyond that is unknown. And the playing-time gap I flagged last week, still true: **1,310 of 1,959 players, 67 percent**, have no matching on-field performance record. I confirmed that's a real coverage gap, not a bug, last week."

IF ASKED WHETHER THESE CAVEATS AFFECT THE VALUE-BY-AGE CHART

No, and it's worth being precise about that. Both caveats come from the injury and performance tables. The value-by-age chart uses only the value history table, so it isn't affected by either one.

12

STATUS CHECKLIST

Raw Analysis, 100% Complete

SAY

"This stage is complete. I connected directly to the dbt marts. I ran sanity checks before drawing any conclusions. I confirmed market value peaks in the 24 to 26 age range. I found 40 percent of players carry no recorded injury history. I documented both data quality caveats honestly. And I'm ready to bring these patterns into the fine-tune stage."

13

CLOSING

Next Update & Thank You

SAY

"My next update is the fine-tune stage: a Prophet value forecast, an injury-risk score, and a combined recruitment signal, all building on what I showed today. Thank you for your time, I'm happy to take questions now."

WRAPPING UP

MOVING INTO DISCUSSION

After slide 11, pause, take a breath, and let the room lead. It's fine to say "let me check my notes" before answering something specific, that reads as careful, not unprepared.

If a question reaches ahead into forecasting or risk scoring, it's fine to say plainly that those are the fine-tune stage, coming at the next update, today was about the raw patterns underneath them.

Quick facts, in case you're asked

A fast lookup, so you don't have to hold every number in your head.

99.98% of injury dates parsed cleanly, 107,951 of 107,971

18,824 of 18,825 shared player IDs overlap between the values and injuries sources

1,959 players analyzed, same combined dataset as last week

32,544 valuation snapshots, 2003 to 2026

24-26 age band where median market value peaks (5.0M EUR, mean 12.4M)

40% of players, 783 of 1,959, show zero recorded injuries

12 players sit exactly at the injury-count cap of 60 (source pagination limit)

67% of players, 1,310 of 1,959, have no matching on-field performance record

**Python, pandas,
DuckDB, dbt**

tools behind both the analysis and the
native chart build

Raw Analysis · Weekly Stakeholder Update · September 10, 2026